

Page Table

- Consider a 32-bit logical address space
- Page size: 4 KB (2^{12})
- Number of pages = $2^{32} / 2^{12} = 2^{20}$
- Page table would have 2^{20} entries
- If each entry is 4 bytes:
 - Space for page table for each process = $2^{20} * 4 \text{ bytes} = 4 \text{ MB}$
- Do not want to allocate space contiguously in main memory
- Solutions:
 - Hierarchical Paging
 - Hashed Page Tables
 - Inverted Page Tables

Logical address space

Page 0
Page 1
Page 2
Page $2^{20} - 1$

Page table

0	
1	
2	
3	<i>f</i>


Hierarchical Page Tables

- Page the page table
 - Multi-level page table
- 32-bit logical address, 4K page size, is divided into:
 - Page number: 20 bits
 - Page offset: 12 bits




Hierarchical Page Tables


- Page size = 4 KB, size of each entry = 4 bytes
 - No. of entries of page table in one page = $4 \text{ KB} / 4\text{B} = 2^{10}$
- To page the page table, page number further divided into:
 - 10-bit page number
 - 10-bit page offset
- Thus, a logical address is as follows:

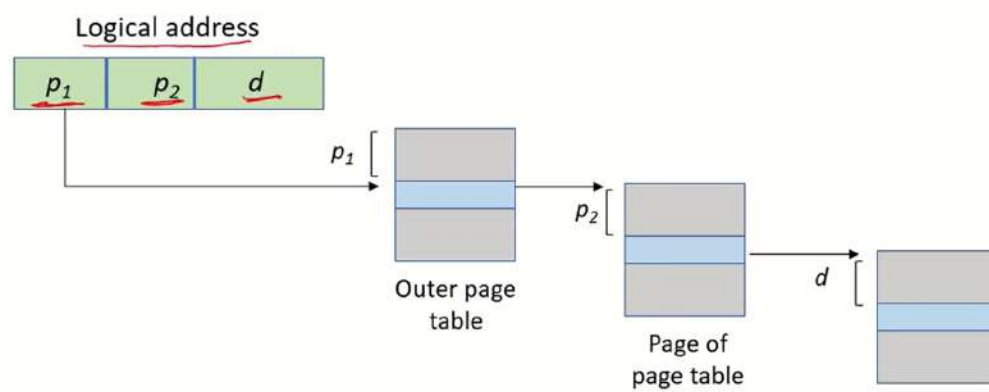


- p_1 is an index into outer page table
- p_2 is displacement within page of inner page table

[illegible]

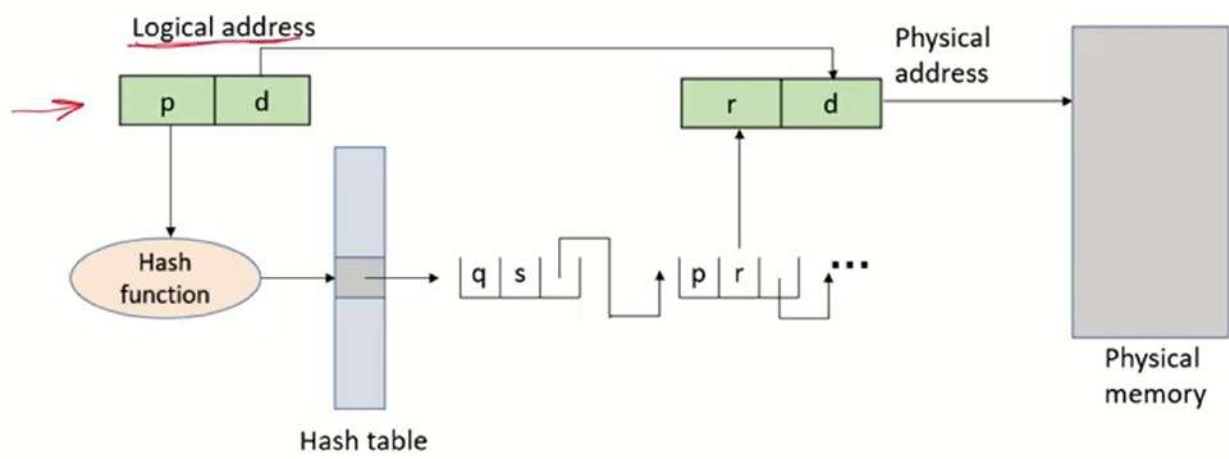


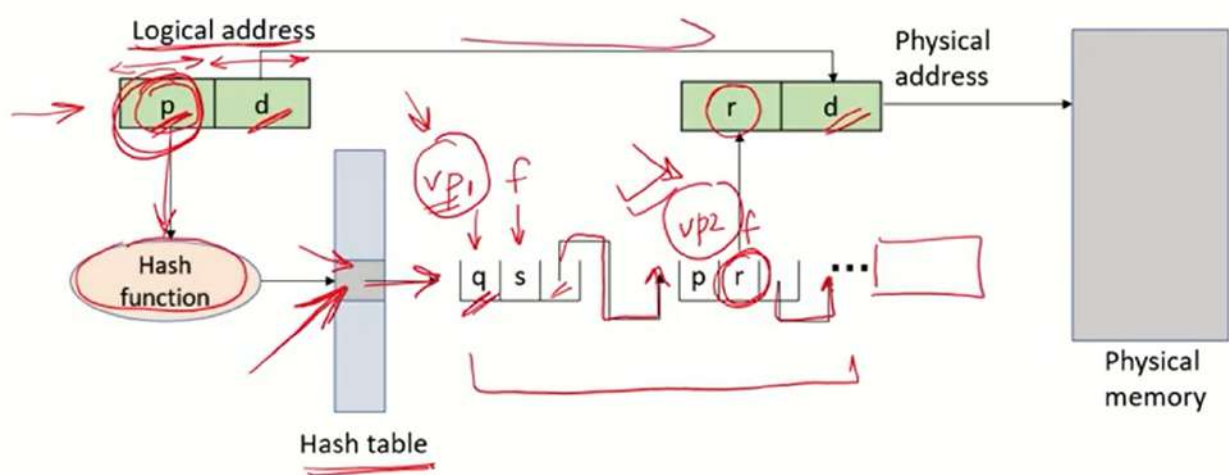




Hashed Page Tables

- Common in address spaces > 32 bits
- Virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to same location
- Each element contains:
 - Virtual page number
 - Value of mapped page frame
 - Pointer to next element
- Virtual page numbers are compared in this chain to look for a match
 - If a match is found, corresponding physical frame is extracted





Inverted Page Table

- Each process has its own page table with entry for each logical page
- Rather than each process having a page table and keeping track of all possible logical pages, track page frames
 - One entry for each frame of memory
 - Entry consists of virtual address of page stored in frame along with information about process that owns that page
- Only one page table in the system

Inverted Page Table

0	4
1	0
2	1
3	
4	
5	
6	
7	

Page Table for P1



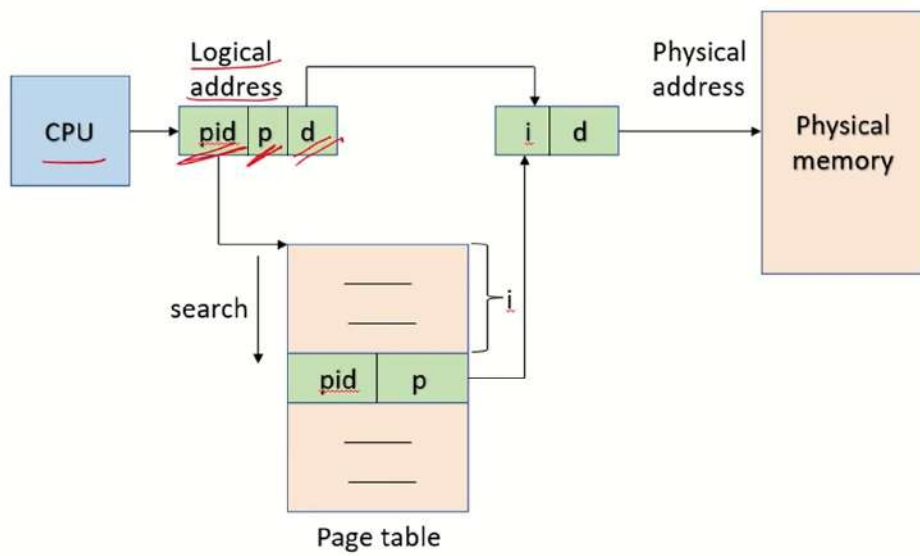
0	2
1	3
2	5
3	
4	
5	
6	
7	

Page Table for P2



	<u>pid</u>	Page no.
0	1	1
1	1	2
2	2	0
3	2	1
4	1	0
5	2	2

Inverted Page Table



Inverted Page Table

- Decreases memory needed to store each page table
 - But increases time needed to search table when a page reference occurs
- Use hash table to limit search to one — or at most a few — page-table entries
 - TLB can accelerate access

Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought **back** into memory for continued execution
 - Improves main memory utilization
 - Swapped-out processes stored in **swap space**
- **Backing store** – fast disk:
 - Large enough to accommodate copies of all memory images for all users
 - Must provide direct access to these memory images
- **Swap out, swap in (Roll out, roll in):**
 - Swapping variant used for priority-based scheduling algorithms
 - Lower-priority process swapped out so higher-priority process can be loaded and executed

Swapping


- Swapping makes possible for total physical address space of all processes to exceed real physical memory of system:
 - Increasing degree of multiprogramming in a system
- Does swapped out process need to swap back in to same physical addresses?
 - Depends on address binding method
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Swapping with Paging

- Standard swapping ~~not used~~ in modern operating systems
 - Swap only when free memory extremely low
- Most systems use a variation of swapping:
 - Pages of a process—rather than an entire process—can be swapped
 - A **page out** operation moves a page from memory to the backing store
 - Reverse process is known as a **page in**


Swapping with Paging

- Standard swapping not used in modern operating systems
 - Swap only when free memory extremely low
- Most systems use a variation of swapping:
 - Pages of a process—rather than an entire process—can be swapped
 - A page out operation moves a page from memory to the backing store
 - Reverse process is known as a page in




Virtual memory

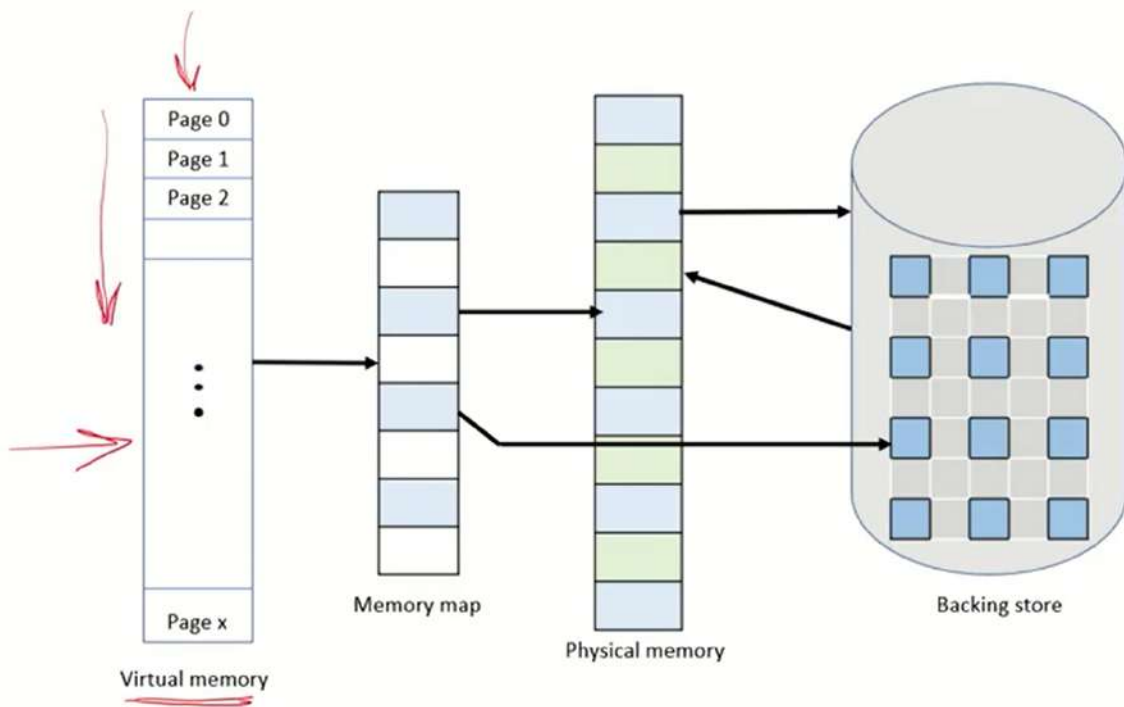
- Instructions must be in physical memory to be executed
- It limits size of a program to size of physical memory
- In many cases, entire program is not needed
 - Codes to handle error conditions
 - Static memory allocation to arrays
 - Certain features of programs
- Entire program may not be needed in memory at same time



Virtual memory

- Ability to execute a program only partially in memory can have many benefits:
 - Program no longer constrained by amount of available physical memory
 - If each program takes less physical memory, more programs can run at the same time - increase in CPU utilization and throughput
 - Less I/O needed to load or swap portions of programs into memory
- **Virtual memory** involves separation of logical memory as perceived by developers from physical memory
 - Allows a large virtual memory to be provided for programmers even if smaller physical memory is available





Virtual memory

- Separation of user logical memory from physical memory
 - Only part of program needs to be in memory for execution
 - Logical address space can be much larger than physical address space
 - Allows address space to be shared by several processes
 - More programs running concurrently
- Logical view of how process is stored in memory
 - Contiguous addresses, consisting of pages
- Physical memory organized in page frames
 - MMU must map logical to physical

Demand Paging

- Can bring entire process into memory at load time
 - May not initially need entire program in memory
- **OR** bring a page into memory only when it is needed (**demanded**)
 - Less I/O needed, no unnecessary I/O
 - Less memory needed
 - Faster response
 - More users
- Page is needed \Rightarrow reference to it
 - Invalid reference \Rightarrow abort
 - Not-in-memory \Rightarrow bring to memory

Demand Paging

- Basic concept - load a page in memory only when it is needed
 - While a process is executing - some pages in memory, and some in secondary storage
 - Need hardware support to distinguish between the two
- With each page table entry, **valid-invalid bit** is associated
 - **v**: in-memory – **memory resident**
 - **i**: not-in-memory
 - Initially valid-invalid bit set to **i** on all entries
- During MMU address translation, if valid-invalid bit is **i**, this results in **page fault**

	Frame #	v/i bit
0		
1	5	v
2	3	v
3	10	v
4		i
	...	
		i
		i

Page table

Demand Paging

- Basic concept - load a page in memory only when it is needed
 - While a process is executing - some pages in memory, and some in secondary storage
 - Need hardware support to distinguish between the two
- With each page table entry, **valid-invalid bit** is associated
 - **v**: in-memory – **memory resident**
 - **i**: not-in-memory
 - Initially valid-invalid bit set to **i** on all entries
- During MMU address translation, if valid-invalid bit is **i**, this results in **page fault**

	Frame #	v/i bit
0		
1	5	v
2	3	v
3	10	v
4		i
	...	
		i
		i

Page table

Demand Paging

- Basic concept - load a page in memory only when it is needed
 - While a process is executing - some pages in memory, and some in secondary storage
 - Need hardware support to distinguish between the two
- With each page table entry, **valid-invalid bit** is associated
 - **v**: in-memory – **memory resident**
 - **i**: not-in-memory
 - Initially valid-invalid bit set to **i** on all entries
- During MMU address translation, if valid-invalid bit is **i**, this results in **page fault**

	Frame #	v/i bit
0		
1	5	v
2	3	v
3	10	v
4		i
	...	
		i
		i

Page table

Page

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H

Logical memory




Frame v/i bit

0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

Page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

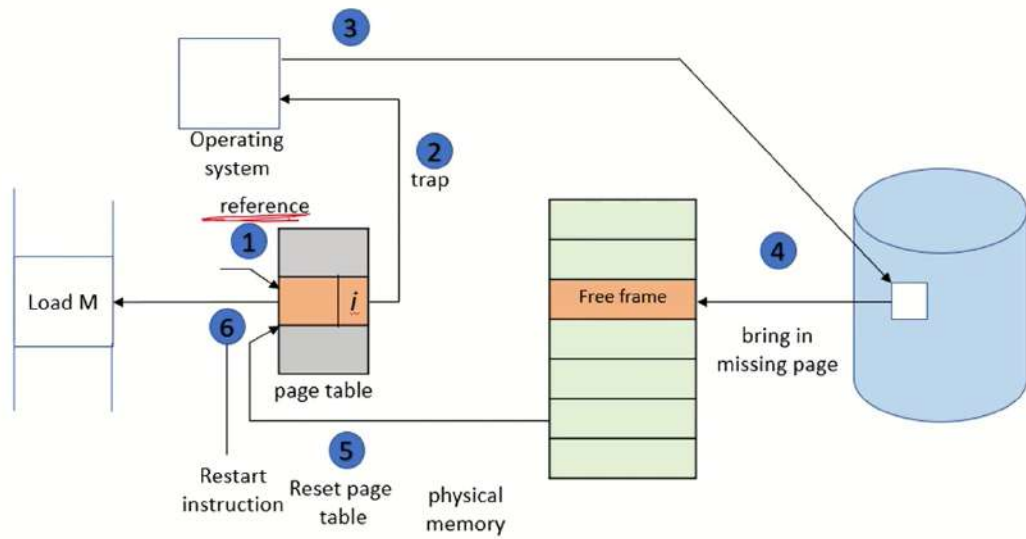
Physical memory



Backing store

Handling Page Faults

1. If there is a reference to a page:
 - Check whether reference is valid or an invalid memory access
 - If invalid – terminate process
- If valid and not in memory - Page fault
 1. Find free frame
 2. Swap page into frame via scheduled disk operation
 3. Reset tables to indicate page now in memory
 - Set validation bit = v
 4. Restart instruction that caused page fault



Free-Frame List

- When page fault occurs – OS must bring desired page from secondary storage into main memory
- Most OS maintain a **free-frame list** -- a pool of free frames for satisfying such requests
 - Typically allocate free frames using **zero-fill-on-demand** technique - content of frames erased before being allocated
- When system starts up, all available memory placed on free-frame list
- As free frames are requested, size of free-frame list shrinks

Performance of Demand Paging

- If no page faults: effective access time = memory access time
- If page fault: first read relevant page from secondary storage and then access desired address
- Let:
 - Memory-access time: ma
 - Probability of a page fault: p ($0 \leq p \leq 1$)
 - Effective access time = $(1 - p) \times ma + p \times \text{page fault service time}$
- Ex.: $ma = 200$ ns, page fault service time = 8 ms
Effective access time = $(1 - p) \times 200 + p \times 8000000 = 200 + 7999800p$
- If one out of 1,000 accesses causes page fault, effective access time is 8.2 μs
 - Slowing down of system by a factor of 40


Copy-on-Write (CoW)



- `fork()` system call creates child process that is duplicate of its parent
- Copy of parent's address space created for child
 - Duplicating pages belonging to parent
- Many child processes invoke `exec()` system call immediately after creation - copying of parent's address space may be unnecessary
- **Copy-on-write** – allows parent and child processes initially to share same pages
 - Shared pages marked as copy-on-write pages
 - Only pages that can be modified marked as copy-on-write
 - Pages that cannot be modified – ex. pages containing executable code - can be shared by parent and child


Copy-on-Write (CoW)

- *fork()* system call creates child process that is duplicate of its parent
- Copy of parent's address space created for child
 - Duplicating pages belonging to parent
- Many child processes invoke *exec()* system call immediately after creation - copying of parent's address space may be unnecessary
- **Copy-on-write** – allows parent and child processes initially to share same pages
 - Shared pages marked as copy-on-write pages
 - Only pages that can be modified marked as copy-on-write
 - Pages that cannot be modified – ex. pages containing executable code - can be shared by parent and child




Copy-on-Write (CoW)


- If either process writes to a shared page:
 - A copy of shared page is created
 - Unmodified pages shared by parent and child processes
- Allows more efficient process creation as only modified pages are copied
- Free pages are allocated from a **pool** of **zero-fill-on-demand** pages
 - Pool should always have free frames for fast demand page execution



Page Replacement

- When a **page fault** occurs, OS must bring desired page from secondary storage into main memory in a **free frame**
- What happens if there is no free frame?
 - Used up by process pages, kernel, I/O buffers, etc.
- **Page replacement** – find some page in memory, currently not in use and free it






Page Fault Service Routine

1. Find location of desired page on disk
2. Find a free frame:
 - If free frame available, use it
 - If no free frame: use page replacement algorithm to select **victim frame**
3. Bring desired page into (newly) free frame
 - Update page and frame tables
4. Continue process by restarting instruction that caused the trap

If no free frame: 2 page transfers for page fault

- Increases effective access time
- Can reduce overhead by using a **dirty** bit (or modify bit)



Dirty Bit

- Each page or frame has a modify bit associated with it
- Modify bit is set whenever page is written into
 - Indicates page has been modified.
- When page selected for replacement, modify bit is examined
- If bit is set - page has been modified while in memory
 - In this case, write page to storage
- If modify bit is not set - page has not been modified
 - Need not write memory page to storage: it is already there


Frame Valid-invalid
bit

<i>o</i>	<i>i</i>
<i>f</i>	<i>v</i>

Page table

2
Change to
invalid

4
Reset page
table for
new page




Physical memory

Swap out
Victim page

1

3
Swap desired
Page in



Backing store

Dirty Bit

- Technique also applies to read-only pages (ex., pages of binary code)
 - Such pages cannot be modified
 - Can be discarded when desired
- Significantly reduces time required to service a page fault
 - Reduces I/O time by one-half if page has not been modified
- many different page-replacement algorithms
- How do we select a particular replacement algorithm?
 - Evaluate an algorithm by running it on a particular string of memory references and computing number of page faults
 - String of memory references is called a **reference string**